



RRC Summer school

Feature matching

Topics of discussion

1.) Feature matching:

- Intro to feature matching and applications.
- Traditional feature matching algorithms.
- Techniques for computing correspondences.

2.) Homography

3.) Image stitching.

What is feature matching?

Feature: Keypoint + descriptor (1D vector in most cases).

In below image, colored lines are correspondences (there may be outliers).

Image 1

Image 2



Key properties of interest points:

- a.) Translation Invariance.
- b.) Rotation invariance.
- c.) Scale invariance
- d.) Robust to changes in illumination.
- e.) Repeatability



Image courtesy: https://www.cc.gatech.edu/classes/AY2016/cs4476_fall/results/proj2/html/zsun311/index.html

What's the use?

a.) Image stitching.



Image stitching

b.) SfM/SLAM.



Building Rome in a day

More applications:

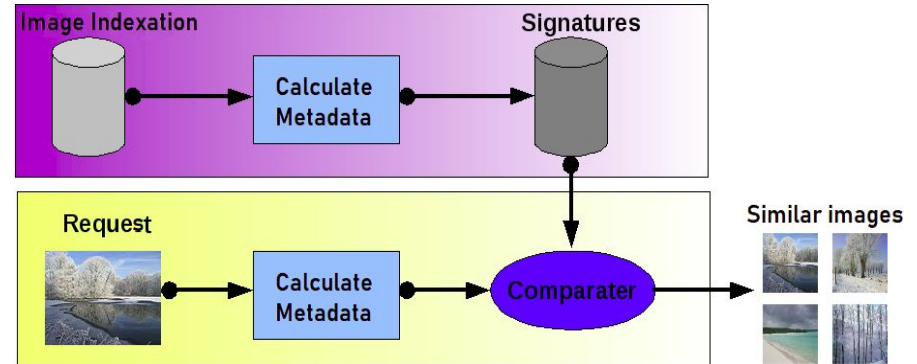
c.) Object tracking.

Remember this match? :)



d.) Image retrieval.

Given query image, return most similar images from DB.



So on.....

Image source: Wikipedia

Famous traditional image descriptors

a.) Harris corner detector

b.) SIFT, SURF

c.) FAST


d.) BRIEF

e.) ORB

So on

Harris corner detection

- Key idea: Sliding window in any direction around the corner results in huge variation in intensity values.



- Consider a pixel of interest:

$$E(u, v) = \sum_{x, y} \underbrace{w(x, y)}_{\text{window function}} \underbrace{[I(x + u, y + v) - I(x, y)]}_{\text{shifted intensity} - \text{intensity}}^2$$

(u, v) is shift in pixels along x and y directions.

Harris corner detection ...

$$E(u, v) = \sum_{x, y} [\underbrace{I(x+u, y+v)}_{\text{shifted intensity}} - \underbrace{I(x, y)}_{\text{intensity}}]^2$$

$$E(u, v) = \sum_{x, y} [\underbrace{I(\cancel{x}, y) + uI_x + vI_y}_{\text{shifted intensity}} - \underbrace{I(\cancel{x}, \cancel{y})}_{\text{intensity}}]^2$$

Taylor Series

$$E(u, v) = \sum_{x, y} [uI_x + vI_y]^2$$

$$E(u, v) = \sum_{x, y} \left[(u \quad v) \begin{pmatrix} I_x \\ I_y \end{pmatrix} \right]^2$$

$$E(u, v) = \sum_{x, y} (u \quad v) \begin{pmatrix} I_x \\ I_y \end{pmatrix} (I_x \quad I_y) \begin{pmatrix} u \\ v \end{pmatrix}$$

$$E(u, v) = (u \quad v) \left[\sum_{x, y} \begin{pmatrix} I_x \\ I_y \end{pmatrix} (I_x \quad I_y) \right] \begin{pmatrix} u \\ v \end{pmatrix}$$



$$M = \sum_{x, y} \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

$$E(u, v) = (u \quad v) M \begin{pmatrix} u \\ v \end{pmatrix}$$

Harris corner detection...

- $E(u, v)$ to be maximized at interest point. This is eqn of ellipse. And shape of ellipse is determined by eigenvalues of M .

Classification of image points using eigenvalues of M :




Should we even compute Eigenvalues? Not efficient to do for many pixels

Harris response value


Denoted by R

$$R = \lambda_1 \lambda_2 - k(\lambda_1 + \lambda_2)^2$$

- Get R for all pixels and select corners if they cross a threshold value.
- Which properties violate?




- Any drawbacks with Harris corner detection?



SIFT...

- Step 1 : Scale space extrema detection (keypoint detection)



First octave

Second octave

Third octave

Fourth octave

Gaussian

In each octave, blurring is increased by a factor of k .

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$


$$B = I * G(k*\sigma)$$

Ref: http://www.cs.cmu.edu/~16385/s17/Slides/7.3_SIFT_Detector_and_Descriptor.pdf

SIFT...


- Approximate LoG using scale space

$$\begin{aligned}\frac{\partial G}{\partial \sigma} &= \sigma \nabla^2 G. \\ \sigma \nabla^2 G &= \frac{\partial G}{\partial \sigma} \approx \frac{G(x, y, k\sigma) - G(x, y, \sigma)}{k\sigma - \sigma} \\ G(x, y, k\sigma) - G(x, y, \sigma) &\approx (k-1)\sigma^2 \nabla^2 G.\end{aligned}$$



SIFT...

- Find the extrema using DoG images



These are not exact optima in LoG, needs to be refined.

Ref: <https://www.cs.ubc.ca/~lowe/papers/ijcv04.pdf>

SIFT...

- **Step 2 : Local Extrema detection**

$$D(\mathbf{x}) = D + \frac{\partial D^T}{\partial \mathbf{x}} \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x}$$



$$\hat{\mathbf{x}} = -\frac{\partial^2 D^{-1}}{\partial \mathbf{x}^2} \frac{\partial D}{\partial \mathbf{x}}.$$



SIFT...

- **Step 3 : Estimate orientation of keypoint**

Why are we estimating direction of keypoint?

$$m(x, y) = \sqrt{(L(x+1, y) - L(x-1, y))^2 + (L(x, y+1) - L(x, y-1))^2}$$


$$\theta(x, y) = \tan^{-1}((L(x, y+1) - L(x, y-1)) / (L(x+1, y) - L(x-1, y)))$$



Ref: <https://www.cs.ubc.ca/~lowe/papers/ijcv04.pdf>

SIFT...

- Step 4 : Compute descriptor




Divide into 8 bins (Total descriptor len = $4 \times 4 \times 8$)

Ref: <https://www.cs.ubc.ca/~lowe/papers/ijcv04.pdf>

FAST (Features from accelerated segment test)

- Works in real time (useful in applications like SLAM, tracking etc).
- Key terms in algo: Thresh (T), bresenham circle (rad 3).
- **Threshold criterion:** N(12) out of 16 pixel intensities should lie outside $[I_p - T, I_p + T]$.
- **Qualifying test:** At least 3 out of 4 pixels (1, 5, 8, 13) should satisfy threshold criterion.
- If pixel p passes Qualifying test, go for Threshold criterion.



Ref: https://homepages.inf.ed.ac.uk/rbf/CVonline/LOCAL_COPIES/AV1011/AV1FeaturefromAcceleratedSegmentTest.pdf

Ref 2 : <https://ieeexplore.ieee.org/document/4674368>

BRIEF (Binary Robust Independent Elementary Features)

- Binary descriptor unlike prev ones. So, matching correspondences is very fast.

Steps:

1. Smooth image using Gaussian kernel (9*9). Why?
2. Get the keypoints using FAST or any other algo.
3. Compute descriptors for each kp:
 - Consider patch of $S \times S$ around kp.
 - Select n_d (x,y) pixel pair locations in patch. How do we do this?
 - Perform test T on each pixel pair.


$$\tau(\mathbf{p}; \mathbf{x}, \mathbf{y}) := \begin{cases} 1 & \text{if } \mathbf{p}(\mathbf{x}) < \mathbf{p}(\mathbf{y}) \\ 0 & \text{otherwise} \end{cases},$$

- Descriptor is computed as:

$$f_{n_d}(\mathbf{p}) := \sum_{1 \leq i \leq n_d} 2^{i-1} \tau(\mathbf{p}; \mathbf{x}_i, \mathbf{y}_i) .$$

Note: BRIEF-k (k is no of bytes per descriptor = $n_d/8$)

Ref: https://www.cs.ubc.ca/~lowe/525/papers/calonder_eccv10.pdf




Selecting test locations in BRIEF:

There are 3 main sampling strategies for test locations:

1.) Uniform:

$$(X, Y) \sim \text{i.i.d. Uniform}\left(-\frac{S}{2}, \frac{S}{2}\right)$$



2.) Gaussian:

$$(X, Y) \sim \text{i.i.d. Gaussian}\left(0, \frac{1}{25}S^2\right)$$



3.) Gaussian with second pixel centered around first pixel:

$$X \sim \text{i.i.d. Gaussian}\left(0, \frac{1}{25}S^2\right)$$

$$Y \sim \text{i.i.d. Gaussian}\left(x_i, \frac{1}{100}S^2\right)$$




ORB (Oriented FAST and Rotated BRIEF):

- Combines idea of FAST kp detector and BRIEF descriptor.
- Any issues with BRIEF and FAST?

Steps:

1. Detect Feature points/keypoints (based on improved FAST):
 - Apply standard FAST to get feature points and compute Harris response values.
 - Sort the feature points based on response values and pick largest N.
 - ORB uses image pyramid technique (to address scale invariance).



ORB ...

- Intensity Centroid method is used to compute direction of FAST features/kps.

$$m_{pq} = \sum_{x,y} x^p y^q I(x,y), \quad \longrightarrow \quad C = \left(\frac{m_{10}}{m_{00}}, \frac{m_{01}}{m_{00}} \right) \quad \longrightarrow \quad \theta = \text{atan2}(m_{01}, m_{10}),$$

2. Computing descriptors for the oriented FAST keypoints:

- Standard BRIEF is not rotation invariant. So, improved version called steerBRIEF is used.
- We already have theta(orientation) of each kp from improved FAST.


$$R_{\theta} = \begin{pmatrix} \cos \theta & \sin \theta \\ -\sin \theta & \cos \theta \end{pmatrix} \quad Q = \begin{bmatrix} x_1, x_2, \dots, x_N \\ y_1, y_2, \dots, y_N \end{bmatrix} \quad \longrightarrow \quad Q_{\theta} = R_{\theta} Q$$

$$g_{N(p,\theta)} = f_N(p) | (x_i, y_i) \in Q_{\theta}$$


3. Matching/Correspondence estimation: Based on FLANN.

Find Correspondences

- **Brute force:** Every descriptor in first image is matched with every descriptor in second image and we check for the closest one based on some distance metric.



- **FLANN** (Fast Library for Approximate Nearest Neighbors): Builds a KDtree/ Hierarchical k means tree from the descriptors of second image (inverted file index). For a descriptor in first image, we query the tree and find the leaf node which is nearest.



Next topic:

1.) Feature matching


2.) Homography and Image stitching:

- Introduction to Homography.
- Computing H using DLT.
- Application of Homography (Panorama)

3.) Code demo.

Homography





Review of 2D transformations:



Ref: CMU 16-385 course

2D transformations ...


Hierarchy of 2D transformations:

Group	Matrix	Distortion	Invariant properties
Projective 8 dof	$\begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix}$		Concurrency, collinearity, order of contact : intersection (1 pt contact); tangency (2 pt contact); inflections (3 pt contact with line); tangent discontinuities and cusps. cross ratio (ratio of ratio of lengths).
Affine 6 dof	$\begin{bmatrix} a_{11} & a_{12} & t_x \\ a_{21} & a_{22} & t_y \\ 0 & 0 & 1 \end{bmatrix}$		Parallelism, ratio of areas, ratio of lengths on collinear or parallel lines (e.g. midpoints), linear combinations of vectors (e.g. centroids). The line at infinity, l_∞ .
Similarity 4 dof	$\begin{bmatrix} sr_{11} & sr_{12} & t_x \\ sr_{21} & sr_{22} & t_y \\ 0 & 0 & 1 \end{bmatrix}$		Ratio of lengths, angle. The circular points, I, J (see section 2.7.3).
Euclidean 3 dof	$\begin{bmatrix} r_{11} & r_{12} & t_x \\ r_{21} & r_{22} & t_y \\ 0 & 0 & 1 \end{bmatrix}$		Length, area

Ref: MVG book

Homography (projective transformation)

- Maps points from one plane to another plane (common centre of projection).



$$x' = Hx$$

Computing H (DLT algorithm):

- Consider a single correspondence b/w 2 images.



$$x^{(1)} = \begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix}$$

$$x^{(2)} = \begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix}$$

$$x^{(2)} = H x^{(1)} \Rightarrow \begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix} = \begin{pmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & 1 \end{pmatrix} \begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix}$$

$$\Rightarrow x_2 = \frac{h_{11}x_1 + h_{12}y_1 + h_{13}}{h_{31}x_1 + h_{32}y_1 + 1}, \quad y_2 = \frac{h_{21}x_1 + h_{22}y_1 + h_{23}}{h_{31}x_1 + h_{32}y_1 + 1}$$

Arranging in matrix form

$$\begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 & -x_1x_2 & -y_1x_1 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -x_1y_1 & -y_1y_2 \end{bmatrix} \begin{bmatrix} h_{11} \\ h_{12} \\ h_{13} \\ h_{21} \\ h_{22} \\ h_{23} \\ h_{31} \\ h_{32} \end{bmatrix} = \begin{bmatrix} x_2 \\ y_2 \end{bmatrix}$$

A

$$\Rightarrow [AH = b]$$

DLT ..

- We get 2 equations for each correspondence, how many are required to solve for H?




Diagram illustrating the Direct Linear Transform (DLT) process. Two images, I_1 and I_2 , are shown. Four corresponding points are identified and labeled 1, 2, 3, and 4. Arrows indicate the mapping from I_1 to I_2 .

The system of equations for each correspondence is:

$$\begin{aligned} A_1 h &= b_1 \\ A_2 h &= b_2 \\ A_3 h &= b_3 \\ A_4 h &= b_4 \end{aligned}$$

These equations are combined into a matrix system:

$$\begin{bmatrix} A_1 \\ A_2 \\ A_3 \\ A_4 \end{bmatrix} h = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{bmatrix}$$

The dimensions of the matrices are indicated as 8×8 for the coefficient matrix and 8×1 for the vector b .

- Solve above system of linear equations using either least squares / SVD.

RANSAC (Random sample consensus):



We have a problem here ..




How do we address the wrong correspondences?

RANSAC ..

Algo:



Randomly
sample points
and fit model.



Set thresh
and count
inliers.

Repeat the above steps for N times. Choose the model which has highest inliers.

Ref: CMU 16-385 course

Adaptive RANSAC:

- What should N be? (N is no of samples - iterations of outer loop)
 - $N = \infty$, sample_count= 0.
 - While $N > \text{sample_count}$ Repeat
 - Choose a sample and count the number of inliers.
 - Set $\epsilon = 1 - (\text{number of inliers})/(\text{total number of points})$
 - Set N from ϵ and (4.18) with $p = 0.99$.
 - Increment the sample_count by 1.
 - Terminate.

$$N = \log(1 - p) / \log(1 - (1 - \epsilon)^s). \quad (4.18)$$

- Epsilon = prob that random data point is outlier.
- p = prob that all points in atleast one of N samples are inliers (set to 0.99).

Ref: MVG book

Computing H using RANSAC:

Steps:

- Get keypoints and match the descriptors to get correspondences.
- RANSAC outer loop (N samples) :
 1. Choose 4 correspondences.
 2. Compute H.
 3. Find no of inliers. (based on error $e = l_2 \text{ norm}(Hx - x')$)
- Choose H_i with highest inliers.
- Recompute H using the model with highest inliers (can be further refined using LM solver).

Image stitching/Panorama

